CIS 472/572, Winter 2018
Homework 4 (Programming): Linear Models
DUE DATES: Submit on Gradescope by Thursday, February 13th at 11:59pm.

1 Overview

In this assignment, you will implement the perceptron algorithm and logistic regression. The file format is the same as in the previous projects: binary attributes and a binary class label represented as comma-separated values.

Please use the provided templates `perceptron.py` and `lr.py` for the perceptron and logistic regression algorithms, respectively. Fill in your code for the training and prediction functions, `train_perceptron`, `predict_perceptron`, `train_lr`, and `predict_lr`. You may add any additional functions you find helpful.

Your code must be your own. **Undergraduates may complete the assignment in teams of 2.** Graduates must complete the assignment alone.

Once you complete the template code, it can be run from the command line with the following arguments:

```
python perceptron.py --maxiter <iters> --model <model> <train> <test>
python lr.py --eta <eta> --l2 <lambda> --maxiter <iters> --model <model> <train> <test>
```

Where `train` is the name of a file containing training data, `test` contains test data to be labeled, `model` is the (optional) filename where you will save the resulting linear model, and `--maxiter` (also optional) is the maximum number of iterations to run. For logistic regression, `eta` is the learning rate and `lambda` is the scale of an L2 regularizer.

2 Model File Format

The provided template code saves model files in the following format:

```
0.10976
foo -1.2857
bar  0.4811
```

where the first line contains the bias, and each line that follows lists one of the attributes and its corresponding weight. You may find this helpful for inspecting the learned weights.

2.1 Perceptron

`train_perceptron` should run the perceptron algorithm until convergence. This means running until all training examples are classified correctly. Your implementation should also halt after 100 passes through the training data, if it has not yet converged. (You do not need to shuffle the example order, except as possible extra credit.)

`predict_perceptron` should return the activation $wx + b$ for a given example x with a given model (w, b).

2.1.1 Logistic Regression

`train_lr` should run the logistic regression learning algorithm for 100 iterations or until convergence. Each iteration should perform one step of batch gradient descent with the given learning rate – do not use stochastic gradient for this assignment. You may assume the algorithm has converged when

the gradient is small (e.g., its magnitude is less than 0.01). On many datasets, logistic regression will use the full 100 iterations without converging.

Use the $\frac{\lambda}{2} \sum_i w_i^2$ as the regularizer, where λ is the strength of the L2 regularizer. The gradient of the regularizer with respect to each w_j is λw_j. When implementing gradient descent, just add the gradient of the regularizer to the gradient of the logistic loss.

`predict_lr` should return the probability that $y = +1$ for a given example x according to a logistic regression model with the given parameters (w, b).

3   Extra Credit

If you have the time, there are lots of ways you can extend this assignment to get a few points of extra credit. Here are a few suggestions, or you can come up with your own extensions.

For perceptron:

- Perform batch gradient descent, minimizing squared error. Compare the results.

- Multi-class classification: explain how you would approach it, and implement your algorithm.

- Implement and evaluate the effects of weight averaging (averaged perceptron algorithm) and randomizing example order on the convergence rate and accuracy of the learned classifier.

For logistic regression:

- Add an L1 prior. Since the L1 prior isn't differentiable at 0, you'll need to implement a sub-gradient method. Compare the results.

- Add a second-order optimization method, such as Newton's method, conjugate gradient, or memory-limited BFGS.

- Implement stochastic gradient optimization. Graph the error rate on the test set as a function of the number of examples processed. How quickly does it converge?

- Implement multi-class classification with softmax.

Or implement a soft-margin linear SVM and compare it to logistic regression and the perceptron.

These are all optional. If you choose to do the extra credit, please turn in your code and write-up **separately**, as different files in your Gradescope submission. Your extra credit code may follow a different interface.